

# Self driving infrastructure

Xiang Li


[xiang.li@coreos.com](mailto:xiang.li@coreos.com) | Head of distributed system

# Topics

- Cluster management systems
- Today's problems with operating cluster management systems
- A self-driving approach

# Motivation: microservices

- Increased operational cost
  - a lot of components
  - dynamic dependencies
  - fast deployment iteration
- Solution: automation




# Cluster management system

- Automation
  - Scheduling
  - Deployment
  - Healing
  - Discovery/load balancing
  - Scaling

# Scheduling











Scheduler



# Scheduling




Scheduler




# Scheduling

Scheduler



# Discovery

 color=yellow






# Discovery




color=yellow

Select color = yellow



# Load balancing


yellow.mycluster  
Select color = yellow



# Healing

Controller manager


● 5



# Healing

Controller manager


● 5




# Healing

Controller manager

● 5





People love automation!

**kubernetes**



kubernetes

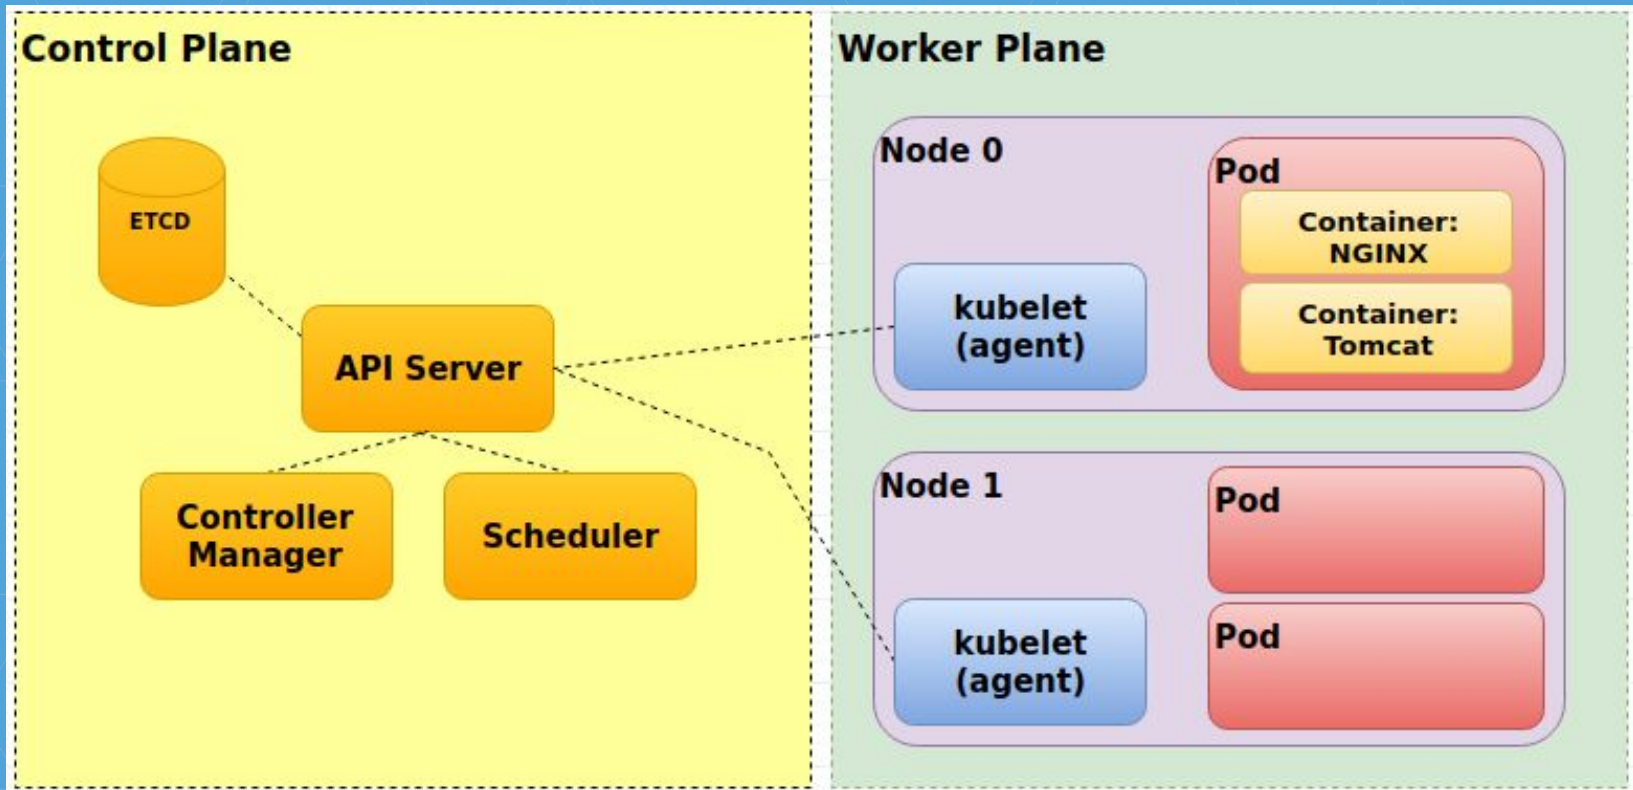


**I hate Kubernetes!**



**I hate to  
OPERATE  
Kubernetes!**

# Kubernetes Architecture



# Operating Kubernetes

- Installation
- Upgrade
- Healing
- Scaling
- Security
- Monitoring
- ...

# Installation

- SSH
- Install kubelet
  - `$pkgmanager install kubelet`
- Install container runtime
  - `$pkgmanager install [docker|rkt]`
- Start kubelet
  - `Systemctl start kubelet`

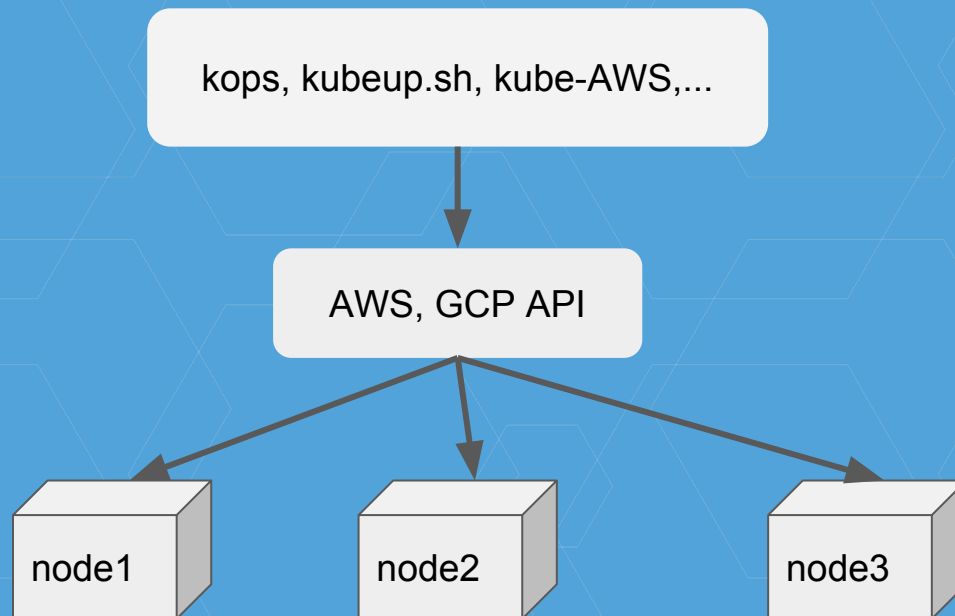
# Installation - master

- SSH
- Install scheduler
- Install controller manager
- Install API server
- Config them correctly
- Start them

# Installation - etcd

- SSH
- Install etcd
- Config them correctly
- Start them

# Installation



# Upgrade

- SSH
- Upgrade container runtime
- Upgrade Kubelet



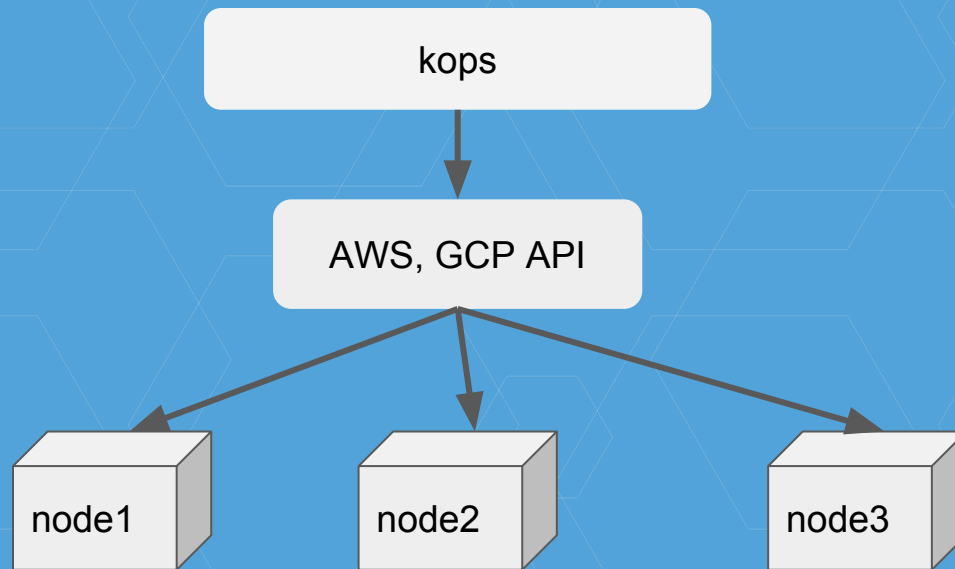
# Upgrade - master

- SSH
- Upgrade master components


# Upgrade - etcd

- SSH
- Upgrade etcd


# Upgrade



# Rollback




# Healing



# Healing


Create node



# Healing

Install

Config



# Problems

A lot of manual/semi-manual work

No standard way to approach all the problems

**do it wrong, lose the cluster!**



# Self hosting

```
// gcc source code
#include <stdio.h>
int main()
{
    compile_c(argv[1]);
}
```



**gcc**



**gcc**

# Self hosting

```
// go lang source code  
package main  
import "os"  
  
func main() {  
    compile_go(os.Args[1:])  
}
```



go



go

# Self hosting



# Self hosting



```
$ uname -s  
minix  
$ gcc linux.c
```


# Self hosting




```
$ uname -s  
minix  
$ gcc linux.c
```



# Self hosting




# Self hosting



```
$ uname -s  
linux  
$ gcc linux.c
```

# Self hosting



```
$ uname -s  
linux  
$ gcc linux.c
```





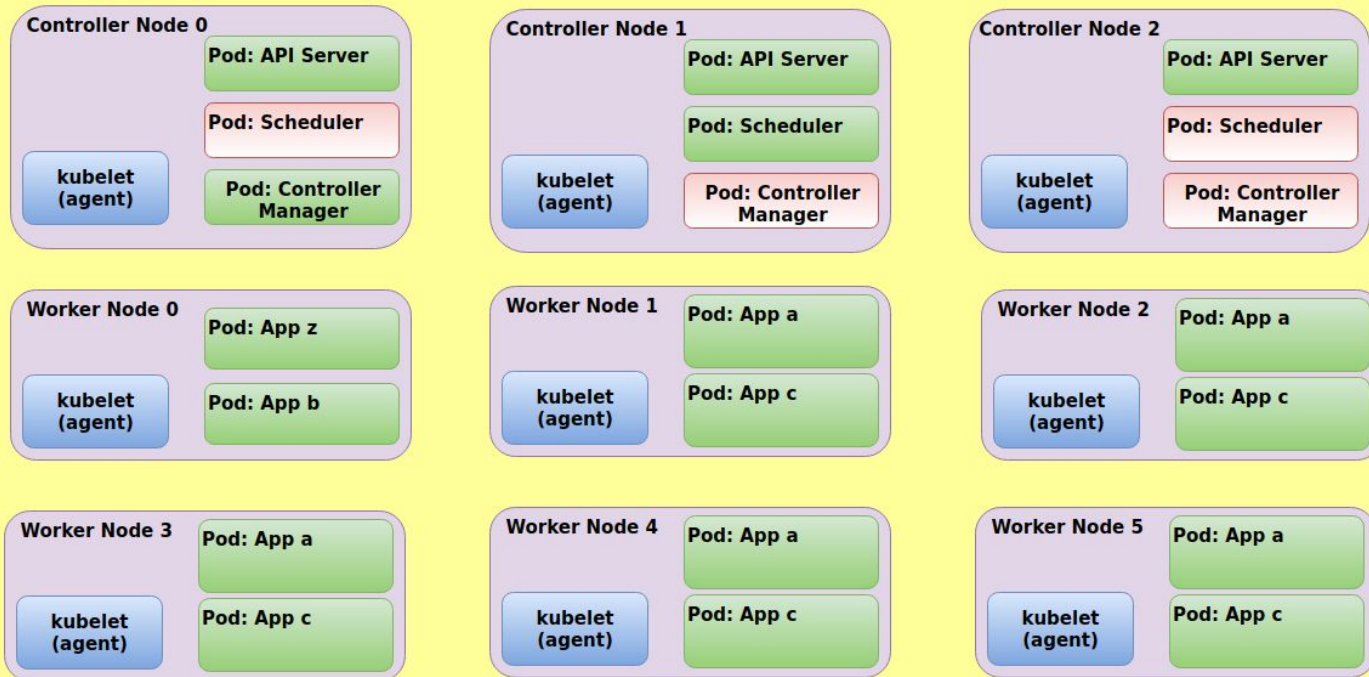
# Self-hosted Kubernetes?

# What is self-hosted Kubernetes?

- Kubernetes manages own core components
- Core components deployed as native API objects

# Self-hosted k8s Architecture

## k8s cluster



# Why Self-host Kubernetes?

- Operational expertise around app management in k8s extends to k8s itself
  - E.g. scaling
- Bootstrapping simplified
- Simply cluster life cycle management
  - E.g. updates
- Upstream improvements in Kubernetes directly translate to improvements in managing Kubernetes

# Simplify Node Bootstrap

On-host requirements become:

- Kubelet
- Container Runtime (docker, rkt, ...)

# Any Distro Node Bootstrap

- Install kubelet
  - `$pkgmanager install kubelet`
- Install container runtime
  - `$pkgmanager install [docker|rkt]`
- Write kubeconfig
  - `scp kubeconfig user@host:/etc/kubernetes/kubeconfig`
- Start kubelet
  - `Systemctl start kubelet`

# Simplify k8s lifecycle management

Manage your cluster with only kubectl

Upgrading a self-hosted Kubernetes cluster:

```
$ kubectl apply -f kube-apiserver.yaml  
$ kubectl apply -f kube-scheduler.yaml  
$ kubectl apply -f kube-controller-manager.yaml  
$ kubectl apply -f kube-proxy.yaml
```

# Launching a self-hosted cluster

Need an initial control plane to bootstrap a self-hosted cluster

Bootkube:

- Acts as a temporary control plane long enough to be replaced by a self-hosted control plane.
- Run only on very first node, then not needed again.


[github.com/kubernetes-incubator/bootkube](https://github.com/kubernetes-incubator/bootkube)

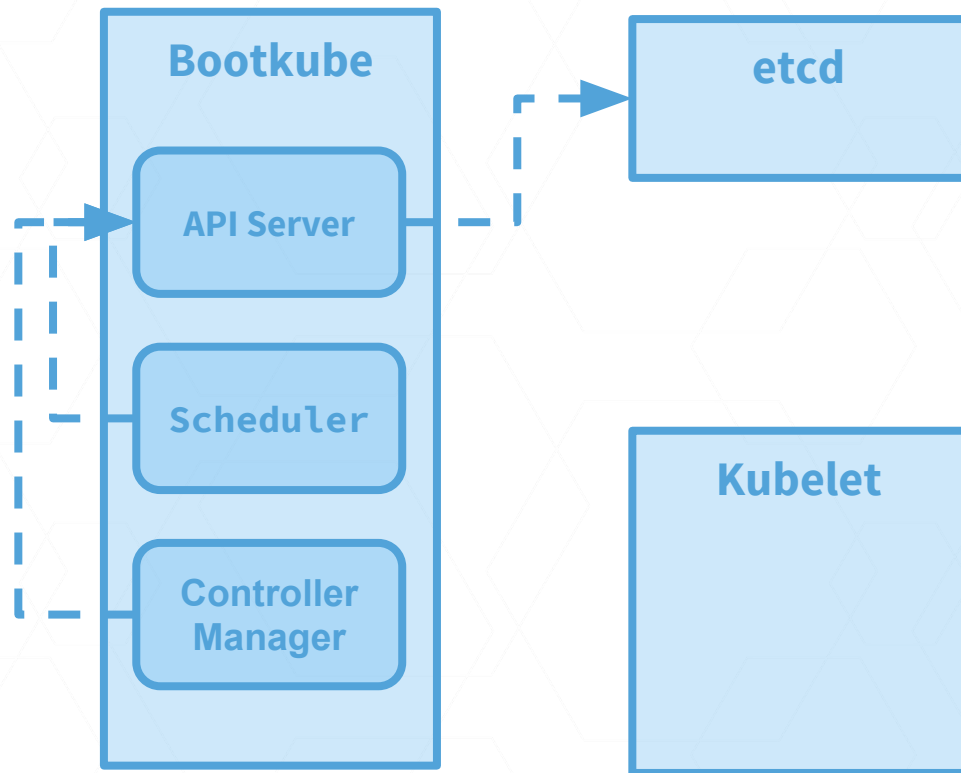



# How Bootkube Works

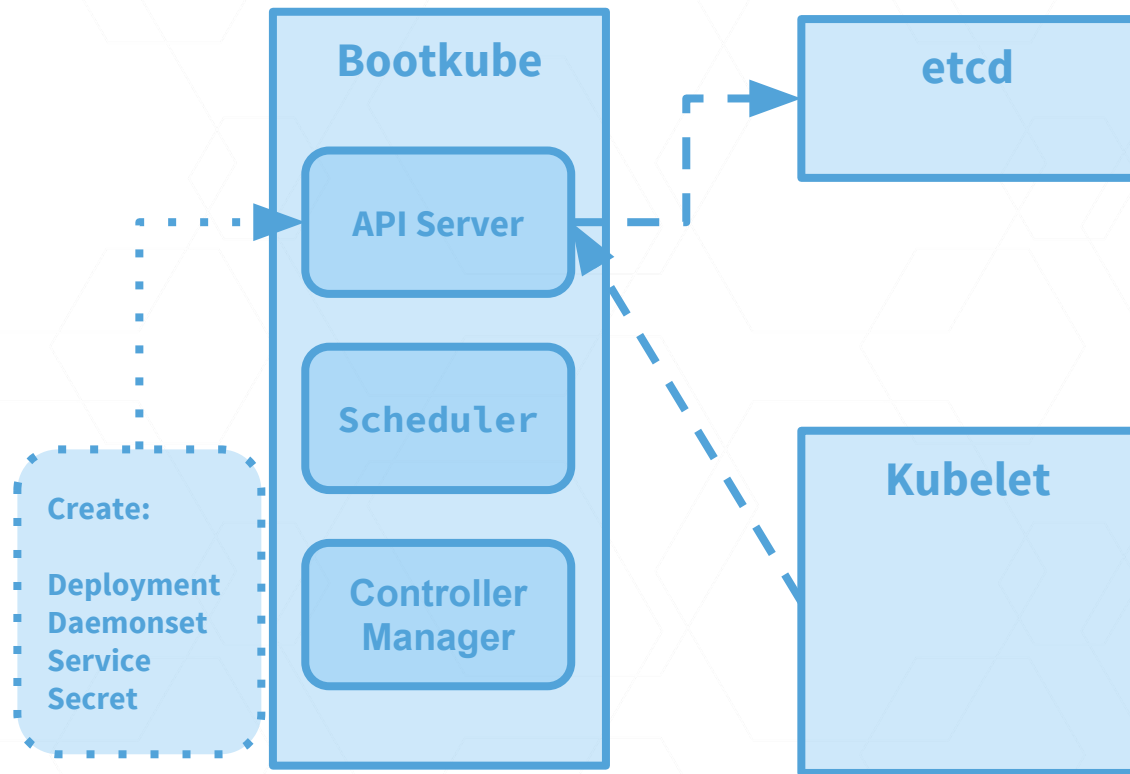
**etcd**


**Kubelet**






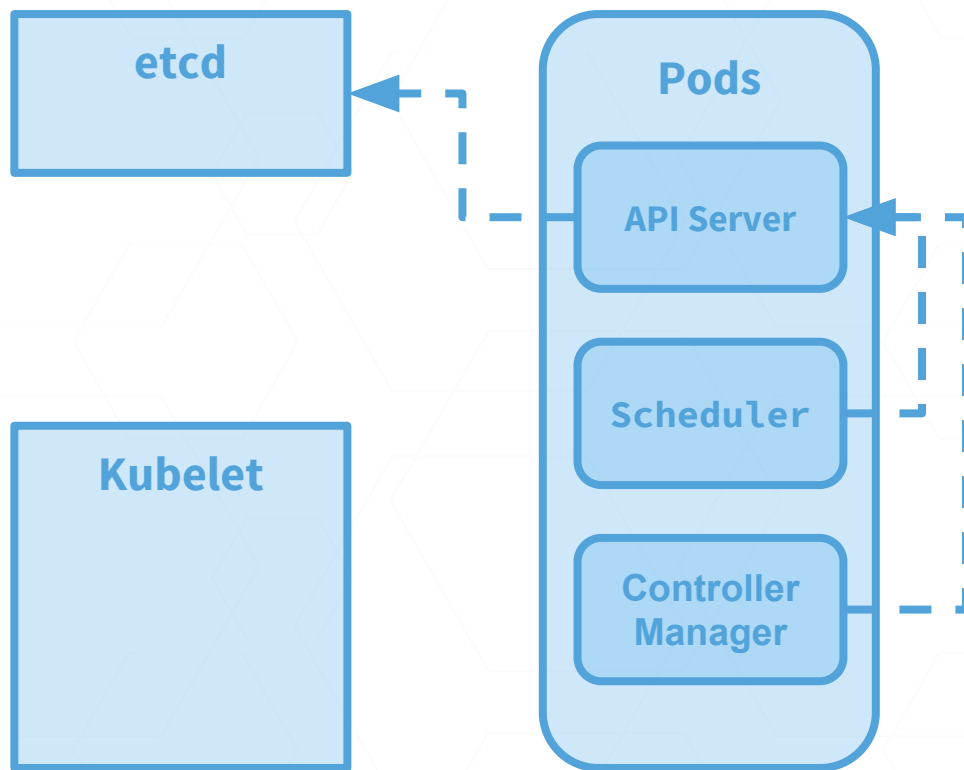















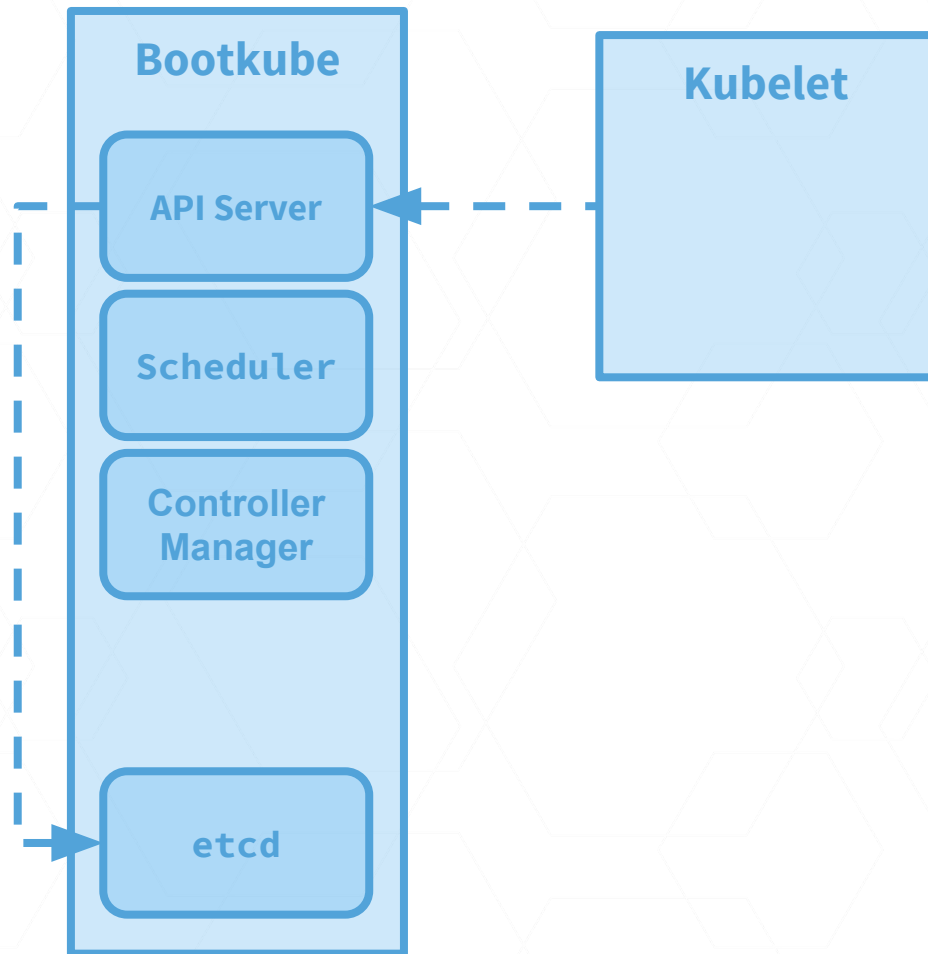
# But wait! There's more!


You can even self-host etcd!


<https://coreos.com/blog/introducing-the-etcd-operator.html>

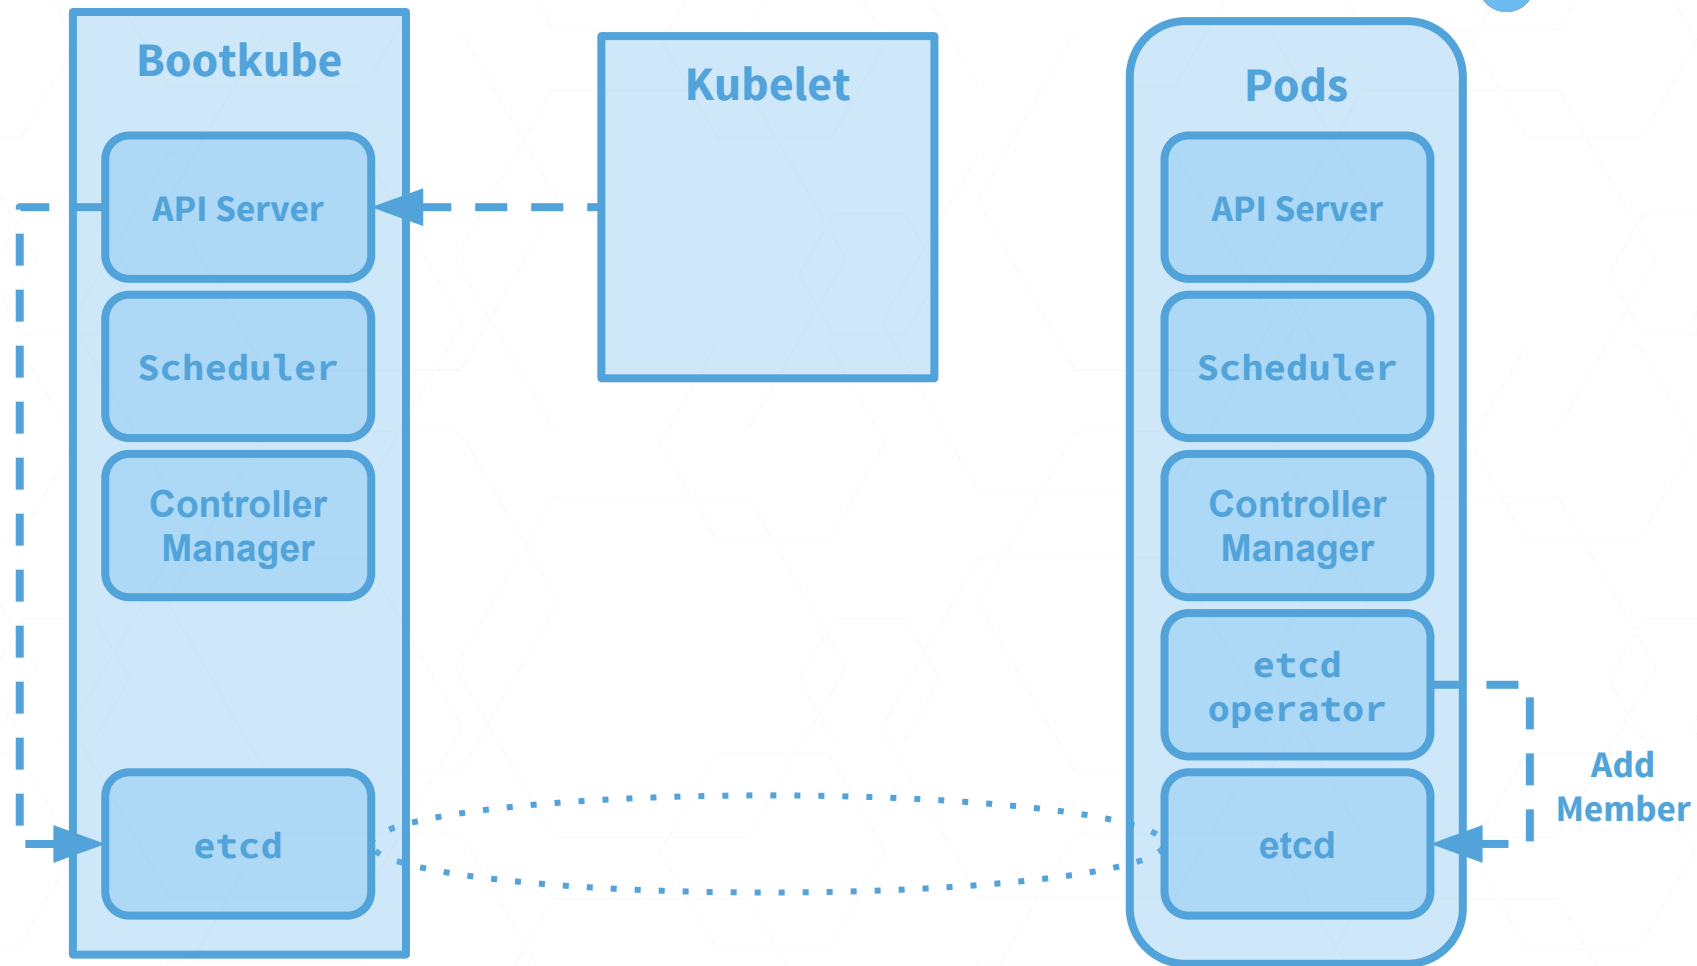
<https://github.com/coreos/etcd-operator>

# How to bootstrap self-hosted etcd




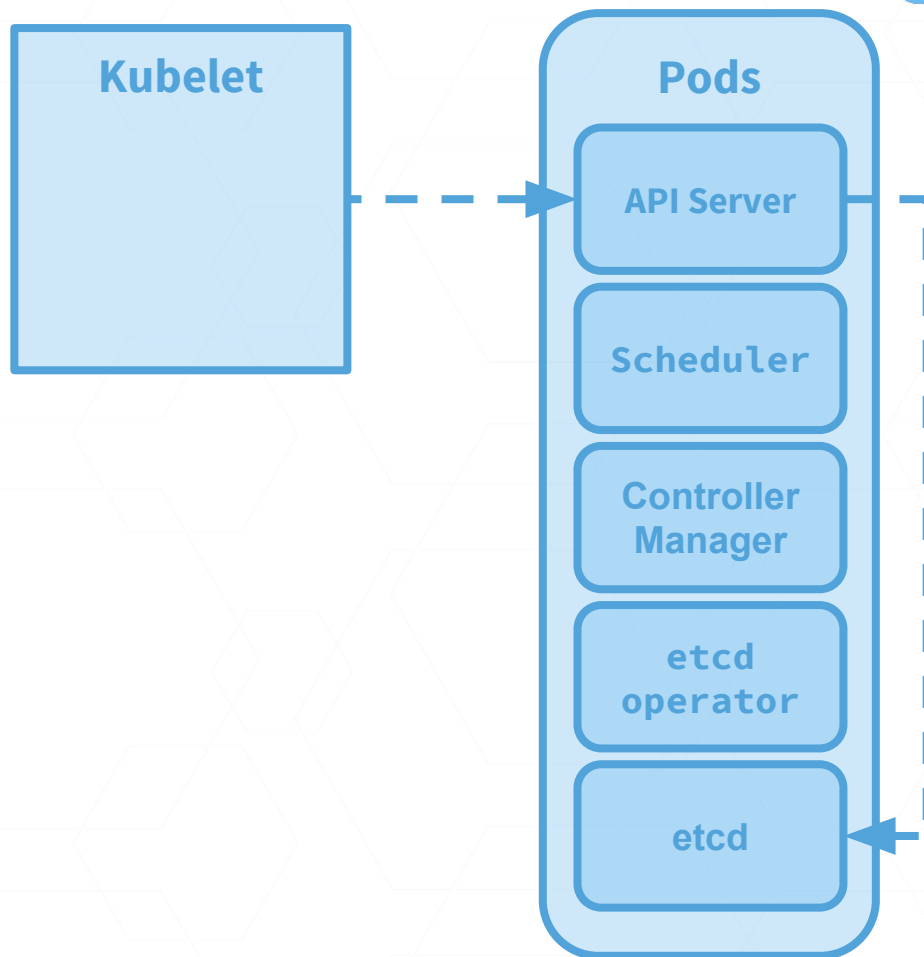












# Disaster Recovery

Node failure in HA deployments (Kubernetes)

Partial loss of control plane components (Kubernetes)

Power cycling the entire control plane (Kubernetes)

Permanent loss of control plane (External tool)

# Disaster Recovery

Permanent loss of control plane

- Similar situation to initial node bootstrap, but utilizing existing etcd state or etcd backup.
- Need to start a temporary replacement api-server
  - Could be binary, static pod, new tool, bootkube, etc.
- Recovery once etcd+api is available can be done via kubectl (as seen previously)


# Self-Driving Kubernetes

# Self driving

- A self-hosted cluster launched via Bootkube
- Upgraded via Kubernetes APIs and an Operator
- Automated by single-button or fully automatic



# Kubernetes Version Operator



Cluster is running v1.4.3 and configured to run v1.4.5

- API Server is v1.4.3
- Scheduler is v1.4.3

Differences from desired config

- API Server should be v1.4.5
- Scheduler should be v1.4.5

How to get there

- Upgrade all API servers Daemons to v1.4.5 safely one-by-one
- Upgrade all Scheduler Deployments to v1.4.5
- Update status to v1.4.5

# The infrastructure

Workload driven

Automation driven

**Easy to manage: self driving approach (Today's topic)**

Security focused





Xiang Li

[xiang.li@coreos.com](mailto:xiang.li@coreos.com)

**Thank you!**